

Top AI Model Providers & Models Released in the Last 12 Months

This document summarizes 10 major AI model providers, their recently released models, release dates, context windows, approximate token pricing, and ideal use cases.

Provider	Model	Release Date	Cutoff Date	Context Window	Approx Token Cost	Best Use Case
OpenAI	GPT-5	Aug 2025	Jun 2025	256K	\$5/\$15 per 1M	Advanced reasoning & coding
Anthropic	Claude 4 Opus	May 2025	Jan 2025	200K	\$15/\$75 per 1M	Long-form analysis
Google	Gemini 2.5 Pro	Mar 2025	Jan 2025	1M	\$3.5/\$10 per 1M	Multimodal reasoning
xAI	Grok 3	Feb 2025	Late 2024	128K	\$5/\$15 per 1M	Real-time knowledge queries
Meta	Llama 4 Maverick	Apr 2025	2024	1M	Open-weight	Self-hosting & research
Mistral	Magistral	Jun 2025	2025	128K	\$2/\$6 per 1M	Efficient enterprise inference
DeepSeek	DeepSeek R1	Jan 2025	2024	128K	\$0.55/\$2.19 per 1M	Reasoning at low cost
Alibaba	Qwen 3 Max	2025	2024	256K	\$1/\$4 per 1M	Multilingual applications
Cohere	Command R+	Aug 2025	2025	128K	\$2.5/\$10 per 1M	Enterprise RAG systems
Baidu	ERNIE 4.5	2025	2024	128K	Enterprise pricing	Mandarin-focused AI

Notes:

- Token pricing shown is approximate input/output cost per 1 million tokens.
- Context window refers to the maximum token length supported by the model.
- Open-weight models allow self-hosting and fine-tuning.
- Reasoning models are optimized for logical problem-solving and multi-step analysis.